

PyForecast Refinement Guide

Fit Quality Analysis and Parameter Learning

January 2026

PyForecast Refinement Guide

Fit Quality Analysis and Parameter Learning

Date: January 2026

Executive Summary

PyForecast includes a refinement module for measuring, logging, analyzing, and improving decline curve fit quality. This system provides:

- * **Hindcast Validation** - Measure actual forecast accuracy through backtesting
- * **Fit Logging** - Persist all fit parameters and metrics for analysis
- * **Residual Analysis** - Detect systematic fit errors (bias, autocorrelation)
- * **Parameter Learning** - Suggest optimal parameters from historical performance
- * **Ground Truth Comparison** - Compare auto-fitted curves against expert ARIES projections

Key Design Principle: All refinement features are disabled by default and observation-only. They don't modify the core fitting behavior—they observe and report on it.

Quick Start

Enable Refinement Features via CLI

```
# Run with hindcast validation
pyforecast process data.csv --hindcast -o output/

# Run with fit logging
pyforecast process data.csv --log-fits -o output/

# Run with residual analysis
pyforecast process data.csv --residuals -o output/

# Enable all refinement features
pyforecast process data.csv --hindcast --log-fits --residuals -o output/
```

Enable via Configuration

```
# pyforecast.yaml
refinement:
  enable_logging: true
  log_storage: sqlite
  log_path: null # Uses ~/.pyforecast/fit_logs.db

  enable_hindcast: true
  hindcast_holdout_months: 6
  min_training_months: 12

  enable_residual_analysis: true

  enable_learning: false
```

Feature 1: Hindcast Validation

Hindcast (backtesting) splits historical data into training and holdout periods, fits on training data, and measures prediction accuracy on holdout.

How It Works

```
Historical Data (24 months)
|-- Training Period (18 months) ? Fit decline curve
\-- Holdout Period (6 months) ? Measure prediction accuracy
```

1. Check data sufficiency: `n_months >= min_training + holdout`
2. Split data: training = first (n - holdout) months
3. Fit on training data only
4. Predict holdout period using fitted model
5. Compute accuracy metrics

Metrics Computed

Metric	Description	Good Value
MAPE	Mean Absolute Percentage Error	< 30%
Correlation	Pearson correlation coefficient	> 0.7
Bias	Systematic over/under prediction	bias < 0.3

CLI Usage

```
# Run hindcast validation
pyforecast process data.csv --hindcast -o output/

# Output includes:
# - output/refinement_report.txt with hindcast summary
# - Per-well MAPE, correlation, and bias
```

Example Output

```
Refinement Analysis:
Hindcast validation: 45 wells
Average MAPE: 18.5%
Good hindcast rate: 78.0%

See output/refinement_report.txt for details
```

Configuration

```
refinement:
  enable_hindcast: true
  hindcast_holdout_months: 6      # Months to hold out
  min_training_months: 12        # Minimum training data
```

Programmatic Usage

```
from pyforecast.refinement import HindcastValidator
from pyforecast.core.fitting import DeclineFitter

validator = HindcastValidator(holdout_months=6, min_training_months=12)
fitter = DeclineFitter()

# Validate a well
result = validator.validate(well, "oil", fitter)

if result:
    print(f"MAPE: {result.mape:.1f}%")
    print(f"Correlation: {result.correlation:.3f}")
    print(f"Good hindcast: {result.is_good_hindcast}")
```

Feature 2: Fit Logging

Fit logging persists all fit parameters, metrics, and diagnostics to SQLite storage for analysis and learning across projects.

Storage Location

Default: `~/.pyforecast/fit_logs.db`

This cumulative database enables learning from all fits over time.

What Gets Logged

Category	Fields
Identification	fit_id, timestamp, well_id, product
Context	basin, formation (if available)
Input	data_points_total, data_points_used, regime_start_idx
Parameters	b_min, b_max, dmin_annual, recency_half_life, regime_threshold
Results	qi, di, b, r_squared, rmse, aic, bic
Residuals	residual_mean, durbin_watson, early_bias, late_bias
Hindcast	hindcast_mape, hindcast_correlation, hindcast_bias

CLI Usage

```
# Enable fit logging during processing
pyforecast process data.csv --log-fits -o output/

# Analyze accumulated fit logs
pyforecast analyze-fits --basin "Permian" -o analysis.csv

# View fit log statistics
pyforecast analyze-fits ~/.pyforecast/fit_logs.db
```

Analyze Fits Command

```
pyforecast analyze-fits [STORAGE_PATH] [OPTIONS]

Options:
  --basin TEXT      Filter by basin name
  --formation TEXT  Filter by formation name
  -p, --product TEXT Filter by product
  -o, --output PATH  Export analysis to CSV
```

Example Output

```
Analyzing fit logs from ~/.pyforecast/fit_logs.db

Fit Log Summary:
  Total fits: 523
  Average R?: 0.892
  Average RMSE: 45.32

Hindcast Performance:
  Fits with hindcast: 485
  Average MAPE: 17.3%
  Good hindcast rate: 82.1%

B-Factor Distribution:
  Mean: 0.524
  Std: 0.312
  Range: [0.010, 1.500]
```

Programmatic Usage

```
from pyforecast.refinement import FitLogger, FitLogAnalyzer

# Log fits during processing
with FitLogger() as logger:
    for well in wells:
        result = fitter.fit(t, q)
        logger.log(result, well, "oil", fitting_config)

# Analyze logged fits
analyzer = FitLogAnalyzer()
summary = analyzer.get_summary(basin="Permian", product="oil")
print(f"Average R?: {summary['avg_r_squared']:.3f}")
```

Feature 3: Residual Analysis

Residual analysis detects systematic fit errors that may not be apparent from R? alone.

Diagnostics Computed

Diagnostic	Description	Ideal Value
Durbin-Watson	Autocorrelation test	~2.0
Early Bias	Error in first half of data	~0%
Late Bias	Error in last half of data	~0%
Autocorr Lag-1	Lag-1 autocorrelation	~0

Durbin-Watson Interpretation

DW Value	Interpretation
< 1.5	Positive autocorrelation (systematic underfitting)
1.5 - 2.5	No significant autocorrelation (good fit)
> 2.5	Negative autocorrelation (overfitting/oscillation)

Validation Issues Generated

Code	Severity	Description
RD001	WARNING	Significant autocorrelation in residuals
RD002	WARNING	Systematic early/late bias pattern
RD003	INFO	Non-zero mean residuals

CLI Usage

```
# Enable residual analysis
pyforecast process data.csv --residuals -o output/

# Output includes residual diagnostics in validation_report.txt
```

Example Output

```
Residual Analysis Summary:  
  Total analyzed: 45  
  With systematic patterns: 8 (17.8%)  
  
Wells with systematic residual patterns:  
  WELL-001/oil: DW=1.23, early_bias=-12.5%, late_bias=8.3%  
  WELL-015/oil: DW=0.98, early_bias=15.2%, late_bias=-5.1%
```

Programmatic Usage

```
from pyforecast.refinement import ResidualAnalyzer  
from pyforecast.refinement.schemas import ResidualDiagnostics  
  
analyzer = ResidualAnalyzer()  
  
# Analyze residuals  
diagnostics = ResidualDiagnostics.compute(actual, predicted)  
  
print(f"Durbin-Watson: {diagnostics.durbin_watson:.2f}")  
print(f"Early bias: {diagnostics.early_bias:.1%}")  
print(f"Late bias: {diagnostics.late_bias:.1%}")  
print(f"Has systematic pattern: {diagnostics.has_systematic_pattern}")  
  
# Get validation issues  
validation_result = analyzer.get_validation_issues(diagnostics, "well_id", "oil")  
for issue in validation_result.issues:  
    print(f"[{issue.code}] {issue.message}")
```

Feature 4: Parameter Learning

Parameter learning analyzes accumulated fit logs to suggest optimal fitting parameters for different basins and formations.

Parameters That Can Be Learned

Parameter	Description
`recency_half_life`	How aggressively to weight recent data
`regime_threshold`	What increase triggers regime detection
`regime_window`	Months of trend data for regime detection
`regime_sustained_months`	How long elevation must persist

Note: Physical parameters (b_min, b_max, dmin) are NOT learned-they're literature-based constraints.

How Learning Works

1. Accumulate fit logs with hindcast data
2. Group by basin/formation (or global)
3. Weight fits by hindcast performance (lower MAPE = higher weight)
4. Derive optimal parameter values
5. Store suggestions in database

Confidence Levels

Sample Count	Confidence
< 20	Low
20 - 99	Medium
? 100	High

CLI Usage

```
# Get parameter suggestions  
pyforecast suggest-params -p oil --basin "Permian"  
  
# Update all suggestions from current fit logs  
pyforecast suggest-params --update  
  
# Export suggestions to CSV  
pyforecast suggest-params -o suggestions.csv
```

Example Output

```
Parameter Suggestion for Permian/oil  
Based on 156 fits (confidence: high)
```

```

Suggested values:
recency_half_life: 10.5
regime_threshold: 0.85
regime_window: 6
regime_sustained_months: 2

Historical performance with these parameters:
Average R2: 0.912
Average hindcast MAPE: 14.2%

```

Programmatic Usage

```

from pyforecast.refinement import ParameterLearner
from pyforecast.refinement.parameter_learning import apply_suggestion

learner = ParameterLearner()

# Get suggestion for a basin
suggestion = learner.suggest(product="oil", basin="Permian")

if suggestion:
    print(f"Suggested half-life: {suggestion.suggested_recency_half_life}")
    print(f"Confidence: {suggestion.confidence}")

    # Apply suggestion to create new config
    new_config = apply_suggestion(suggestion, base_fitting_config)

```

Feature 5: Ground Truth Comparison

Ground truth comparison measures how well pyforecast's auto-fitted decline curves match expert/approved ARIES projections.

Use Case

When you have existing ARIES forecasts created by reservoir engineers, you can compare pyforecast's automated fits against these "ground truth" projections to:

- * Validate that auto-fitting produces reasonable results
- * Identify wells where manual review may be needed
- * Measure overall fitting accuracy across a portfolio

ARIES Expression Format

PyForecast parses ARIES AC_ECONOMIC expression format:

```
"{Qi} X {unit} {Dmin%} {type} B/{b} {Di%}"
```

Example: ``1000 X B/M 6 EXP B/0.50 8.5`` means:

- * $qi = 1000 \text{ bbl/month}$
- * unit = B/M (barrels/month)
- * $dmin = 6\% \text{ annual terminal decline}$
- * type = EXP (exponential)
- * $b = 0.50 \text{ hyperbolic exponent}$
- * $di = 8.5\% \text{ annual initial decline}$

CLI Usage

```

# Run with ground truth comparison
pyforecast process data.csv --ground-truth aries_projections.csv -o output/

# With custom comparison period (default: 60 months)
pyforecast process data.csv --ground-truth aries_projections.csv --gt-months 120 -o output/

# Generate comparison plots for each well
pyforecast process data.csv --ground-truth aries_projections.csv --gt-plots -o output/

# Use lazy loading for large ARIES files (streams instead of loading into memory)
pyforecast process data.csv --ground-truth large_aries.csv --gt-lazy -o output/

# Enable parallel validation for large batches
pyforecast process data.csv --ground-truth aries_projections.csv --gt-workers 4 -o output/

```

Example Output

```

Ground Truth Comparison:
Wells with ARIES data: 45 of 50

```

```
Average MAPE: 12.3%
Average correlation: 0.987
Good match rate: 82.2%
```

```
See output/ground_truth_report.txt for details
```

Metrics Computed

Metric	Description	Good Value
MAPE	Mean Absolute Percentage Error of rates	< 20%
Correlation	Pearson correlation of rate curves	> 0.95
Cumulative Diff	Difference in total production	< 15%
B-Factor Diff	Absolute difference in b-factor	< 0.3

Match Quality Grades

Wells are graded A-D based on overall match quality:

Grade	Criteria
A	Excellent match - all metrics well within thresholds
B	Good match - minor deviations
C	Fair match - some significant differences
D	Poor match - review recommended
X	Insufficient data - MAPE could not be calculated

Quality Threshold (is_good_match)

A well is considered a "good match" when ALL of these criteria are met:

- * MAPE < 20%
- * Correlation > 0.95
- * Cumulative diff < 15%
- * B-factor diff < 0.3

Ground Truth Report

The report (`ground_truth_report.txt`) includes:

```
PyForecast Ground Truth Comparison Report
=====
Summary:
    Total comparisons: 45
    Average MAPE: 12.3%
    Median MAPE: 10.5%
    Average correlation: 0.987
    Average cumulative diff: 8.2%
    Good match rate: 82.2%

Grade Distribution:
    A: 25 (55.6%)
    B: 12 (26.7%)
    C: 5 (11.1%)
    D: 3 (6.7%)

Detailed Results:
-----
WELL-001/oil [A] GOOD
    ARIES:      qi=500.0, di=8.5%/yr, b=0.500
    pyforecast: qi=512.3, di=8.8%/yr, b=0.520
    Differences: qi=+2.5%, di=+3.5%, b=+0.020
    Metrics: MAPE=8.5%, corr=0.992, bias=+3.2%, cum_diff=+5.1%
```

Programmatic Usage

```
from pyforecast import_ import AriesForecastImporter
from pyforecast.refinement import GroundTruthValidator, GroundTruthConfig

# Load ARIES forecasts
importer = AriesForecastImporter()
importer.load("aries_projections.csv")

# Configure comparison
config = GroundTruthConfig(comparison_months=60)
validator = GroundTruthValidator(importer, config)

# Compare a well
result = validator.validate(well, "oil")
```

```

if result:
    print(f"MAPE: {result.mape:.1f}%")
    print(f"Correlation: {result.correlation:.3f}")
    print(f"Good match: {result.is_good_match}")
    print(f"Grade: {result.match_grade}")

```

ARIES CSV File Format

The ARIES CSV file should have columns for well identifier and forecast expressions:

```

PROPNUM,OIL_EXPRESSION,GAS_EXPRESSION
42-001-00001,1000 X B/M 6 HYP B/0.50 8.5,5000 X M/M 6 EXP B/0.10 12
42-001-00002,500 X B/D 6 HYP B/0.75 10,

```

Supported unit codes:

- * `B/M` - barrels/month (oil)
- * `B/D` - barrels/day (oil)
- * `M/M` - mcf/month (gas)
- * `M/D` - mcf/day (gas)

Advanced Features

Rate Validation

Forecast rate arrays are automatically validated for problematic values:

- * **NaN values** are replaced with 0 and logged as warnings
- * **Infinite values** are clipped to 1e9 and logged as warnings
- * **Negative values** are clipped to 0 and logged as warnings

This prevents silent calculation errors in metrics.

MAPE Edge Case Handling

When insufficient valid data points exist (fewer than 3 points above the 0.1 rate threshold), MAPE returns `None` instead of an incorrect value:

- * `mape_valid` property indicates if MAPE was successfully calculated
- * `is_good_match` returns `False` when MAPE is unavailable
- * `match_grade` returns `X` for insufficient data

Identifier Mismatch Logging

The `validate_batch()` method tracks wells that exist in only one dataset:

- * `wells_in_pyf_only`: Wells with pyforecast data but no ARIES data
- * `wells_in_aries_only`: Wells with ARIES data but not in pyforecast batch

This helps diagnose ID normalization issues.

Parse Failure Logging

Unparseable ARIES expressions are logged with warnings and tracked:

```

importer = AriesForecastImporter()
importer.load("aries_projections.csv")

# Access parse failures
for well_id, product, expression in importer.parse_failures:
    print(f"Failed to parse {well_id}/{product}: '{expression}'")

```

Time-Series Export

Forecast arrays are exported to `ground_truth_timeseries.csv` for external visualization:

Column	Description
well_id	Well identifier
product	Product type (oil/gas/water)
month	Forecast month (0-based)
aries_rate	ARIES forecast rate
pyf_rate	pyforecast forecast rate
diff	Rate difference (pyf - aries)
pct_diff	Percentage difference

Comparison Plots

Generate overlay plots with the `--gt-plots` flag:

```
pyforecast process data.csv --ground-truth aries.csv --gt-plots -o output/
```

Plots are saved to `output/ground_truth_plots/` with one PNG per well/product showing:

- * ARIES curve (blue solid line)
- * pyforecast curve (red dashed line)
- * Metrics text box (MAPE, correlation, grade)

Lazy Loading

For large ARIES files (10,000+ wells), use lazy mode to reduce memory usage:

```
pyforecast process data.csv --ground-truth large_aries.csv --gt-lazy -o output/
```

In lazy mode:

- * File is validated and rows counted during `load()`
- * Data is streamed from disk on each `get()` call
- * Uses constant memory regardless of file size
- * Slower per-lookup but suitable for memory-constrained environments

Parallel Validation

For large batches, enable parallel validation:

```
pyforecast process data.csv --ground-truth aries.csv --gt-workers 4 -o output/
```

Uses `ThreadPoolExecutor` with configurable worker count. Default is 1 (sequential).

Regime Detection Calibration

Calibrate regime detection thresholds using known refrac/workover events.

Events File Format

```
well_id,event_date,event_type
42-001-00001,2022-06-15,refrac
42-001-00002,2023-01-20,workover
42-003-00005,2022-09-01,refrac
```

CLI Usage

```
pyforecast calibrate-regime data.csv --events known_events.csv -o calibration.json
```

Example Output

```
Loading production data from data.csv...
Loading known events from known_events.csv...
Found 45 known events

Testing regime detection thresholds...
threshold=0.50: 92.0% detection rate (41/45)
threshold=0.75: 88.0% detection rate (39/45)
threshold=1.00: 78.0% detection rate (35/45)
threshold=1.25: 65.0% detection rate (29/45)
threshold=1.50: 52.0% detection rate (23/45)
threshold=2.00: 35.0% detection rate (16/45)

Recommended threshold: 0.50 (92.0% detection rate)
```

Configuration Reference

Complete Refinement Configuration

```
refinement:
  # Logging settings
  enable_logging: false          # Log fit metadata to storage
  log_storage: sqlite            # Storage type: sqlite or csv
  log_path: null                 # Path to storage (null = ~/.pyforecast/fit_logs.db)

  # Hindcast settings
  enable_hindcast: false         # Run hindcast validation
  hindcast_holdout_months: 6     # Months to hold out
  min_training_months: 12        # Minimum training data required

  # Residual analysis
  enable_residual_analysis: false

  # Regime calibration
  known_events_file: null        # CSV with known events

  # Parameter learning
  enable_learning: false

  # Ground truth comparison
  ground_truth_file: null        # ARIES AC_ECONOMIC CSV for comparison
  ground_truth_months: 60         # Months to compare forecasts
  ground_truth_lazy: false        # Stream file instead of loading into memory
  ground_truth_workers: 1          # Parallel workers (1 = sequential)
```

Ground Truth CLI Flags

Flag	Default	Description
--ground-truth PATH	-	ARIES AC_ECONOMIC CSV file for comparison
--gt-months N	60	Months to compare forecasts
--gt-plots	off	Generate comparison plots for each well
--gt-lazy	off	Stream ARIES file instead of loading into memory
--gt-workers N	1	Number of parallel workers for validation

Refinement Report Format

The refinement report ('refinement_report.txt') is generated in the output directory when refinement features are enabled.

Sample Report

```
PyForecast Refinement Analysis Report
=====
Hindcast Validation Summary:
Wells with hindcast: 45
Average MAPE: 17.3%
Median MAPE: 14.8%
Average correlation: 0.892
Good hindcast rate: 82.2%

Hindcast Details:
-----
WELL-001/oil: MAPE=12.5%, corr=0.945, bias=-3.2% [GOOD]
WELL-002/oil: MAPE=8.3%, corr=0.978, bias=1.5% [GOOD]
WELL-003/oil: MAPE=35.2%, corr=0.654, bias=-18.5% [POOR]
...
Residual Analysis Summary:
Total analyzed: 45
With systematic patterns: 6 (13.3%)

Wells with systematic residual patterns:
-----
WELL-003/oil: DW=1.12, early_bias=-15.2%, late_bias=8.5%
WELL-018/oil: DW=0.95, early_bias=12.8%, late_bias=-4.2%
```

Best Practices

Getting Started with Refinement

- **Start with hindcast validation** - Measure baseline forecast accuracy
- **Enable fit logging** - Accumulate data for learning
- **Run residual analysis** - Identify systematic fit issues
- **After 50+ fits** - Check parameter suggestions

Workflow for Continuous Improvement

```
# Initial processing with all refinement features
pyforecast process data.csv --hindcast --log-fits --residuals -o output/

# Review refinement report
cat output/refinement_report.txt

# After processing many wells, check suggestions
pyforecast suggest-params -p oil --basin "Permian"

# If suggestions look good, update config and re-run
pyforecast process new_data.csv -c updated_config.yaml -o output/
```

Interpreting Poor Hindcast Results

Symptom	Possible Cause	Action
High MAPE	Wrong decline behavior	Check b-factor bounds
Low correlation	Data noise or regime changes	Review regime detection
Large positive bias	Over-prediction	Reduce recency weighting
Large negative bias	Under-prediction	Increase recency weighting

When to Trust Parameter Suggestions

- * [OK] High confidence (100+ samples)
- * [OK] Basin/formation specific
- * [OK] Good hindcast performance (MAPE < 20%)
- * [!] Medium confidence - use with caution
- * [X] Low confidence - gather more data first

API Reference

Data Classes

```
from pyforecast.refinement import (
    FitLogRecord,          # Complete fit metadata
    HindcastResult,        # Hindcast validation results
    ResidualDiagnostics,   # Residual analysis metrics
    GroundTruthResult,     # Ground truth comparison results
)
from pyforecast.refinement.schemas import ParameterSuggestion
```

Validators and Analyzers

```
from pyforecast.refinement import (
    HindcastValidator,      # Hindcast validation
    ResidualAnalyzer,       # Residual analysis
    FitLogger,              # Fit logging
    ParameterLearner,       # Parameter suggestions
    GroundTruthValidator,   # Ground truth comparison
    GroundTruthConfig,      # Ground truth configuration
    GroundTruthSummary,     # Batch validation summary with mismatch info
    summarize_ground_truth_results, # Aggregate statistics
)
```

Plotting

```
from pyforecast.refinement.plotting import (
    plot_ground_truth_comparison, # Single well plot
    plot_all_comparisons,        # Batch plot generation
)
```

ARIES Import

```
from pyforecast import_ import (
    AriesForecastImporter,    # Load ARIES forecast files
    AriesForecastParams,      # Parsed forecast parameters
)
```

Storage

```
from pyforecast.refinement import FitLogStorage
from pyforecast.refinement.storage import get_default_storage_path

# Default path: ~/.pyforecast/fit_logs.db
storage = FitLogStorage()

# Custom path
storage = FitLogStorage("/path/to/custom.db")

# Query fit logs
records = storage.query(basin="Permian", product="oil", min_r_squared=0.8)

# Get statistics
stats = storage.get_statistics(basin="Permian")
```

Troubleshooting

"No fit logs found"

```
Error: No fit logs found at ~/.pyforecast/fit_logs.db
```

****Solution:**** Run `pyforecast process --log-fits` first to create fit logs.

"Need at least 10 fits with hindcast data"

```
No parameter suggestions available.
Need at least 10 fits with hindcast data to generate suggestions.
```

****Solution:**** Run `pyforecast process --log-fits --hindcast` on more wells.

High MAPE on all wells

Possible causes:

1. Data quality issues (check validation report)
2. Wrong b-factor bounds for your play
3. Regime detection not capturing correct decline period

Durbin-Watson consistently low

A DW < 1.5 suggests systematic underfitting. Try:

1. Adjusting b-factor bounds
2. Checking for missed regime changes
3. Reviewing recency weighting (may be too aggressive)